

Kalen Foster

ITAI 1378

June 24, 2026

L05 Reflective Journal

The Chihuahua vs. Muffin lab was an engaging introduction to convolutional neural networks (CNNs) and how they fit into image classification. The whole point of the lab was to train a model that could tell apart pictures of chihuahuas vs muffins, which can look alike in a way, textures, colors, and shapes—at least at first glance. It felt like a continuation of the before workshop, but with a shift from a more traditional neural network to a CNN architecture, and it's much more suited for image work compared to the simpler approach.

The most important concept I learned in this lab was how a CNN is different from a normal feedforward neural network. In the earlier workshop, the neural net treated an image like a flattened list, basically pixel values only, and that wipes out the spatial relationships between pixels. A CNN instead preserves the image structure, by using convolutional layers to scan small regions of the image and pick up on patterns like edges, textures, and shapes. Then pooling layers shrink the size of the feature maps while trying to keep the key information. After that, fully connected layers at the end do the last classification step. So overall, CNNs are especially good for computer vision tasks, because they can learn visual features on their own, rather than depending on manually engineered inputs (LeCun, Bengio, & Hinton, 2015).

This type of image classification model has a lot of practical uses, like medical image analysis, self-driving cars, security setups, and even quality control inside manufacturing. But, also there are ethical things to keep in mind. For example, if the training data is not diverse enough, the model may behave poorly on every day, real-world images. And when it misclassifies something, that can really become harmful in high-stakes environments, like healthcare or law enforcement. So for that reason, image classification systems have to be tested carefully and used with human oversight, whenever it is needed (Goodfellow, Bengio, & Courville, 2016).

Overall, this lab gave me a stronger understanding of how CNNs work and why they're so effective at image classification, but it also makes sense why the architecture matters. It showed me the importance of choosing the right structure for the data and carefully evaluating a model's performance rather than just depending on accuracy. The Chihuahua vs. Muffin example made the whole activity pretty memorable, yet the lessons from the lab still apply to far more serious, real world machine learning issues.

Citations

Goodfellow, I., Bengio, Y., & Courville, A. (2016). *Deep learning*. MIT Press.

LeCun, Y., Bengio, Y., & Hinton, G. (2015). Deep learning. *Nature*, 521(7553), 436–444.